

ML/DL for Everyone Season2

with  TensorFlow

Lab 05-2 Logistic Regression


cost function & optimizer

code: <https://github.com/deeplearningzerotoall/TensorFlow>
slides: <http://bit.ly/2LQMKvk>
lecturer: SuSang Kim (healess@kaist.ac.kr)



Logistic Regression




- What is Logistic Regression?
 - Classification
 - Logistic vs Linear
- How to solve?
 - Hypothesis Representation
 - Sigmoid/Logistic Function
 - Decision Boundary
 - Cost Function
 - Optimizer (Gradient Descent)
- Codes (Eager Execution)
- Summary



Cost Function

the cost function to fit the parameters (θ)

W



Given the training set how to we chose/fit θ ?

$h_{\theta}(x) = y$ then Cost = 0


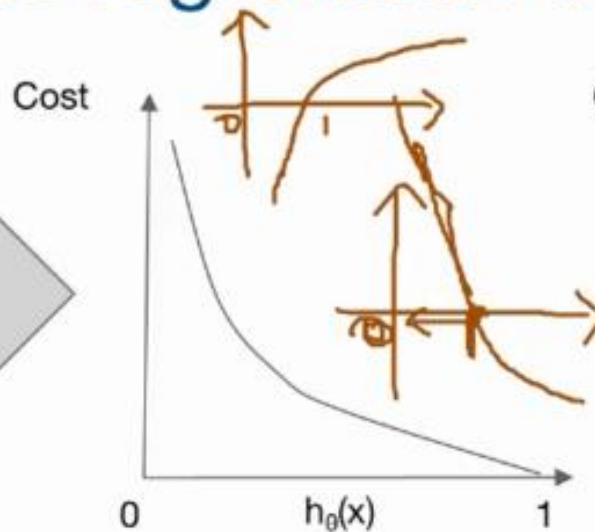


$$\text{cost}(h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1-y) \log(1 - h_{\theta}(x))$$

[Tensorflow Code]

```
def loss_fn(hypothesis, labels):  
    cost = -tf.reduce_mean(labels * tf.log(hypothesis) + (1 - labels) * tf.log(1 - hypothesis))  
    return cost
```

Cost Function

A convex logistic regression cost function



$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_\theta(x^{(i)}) - y^{(i)})^2$$

[Tensorflow Code]

```
cost = -tf.reduce_mean(labels * tf.log(hypothesis) + (1 - labels) * tf.log(1 - hypothesis))
```


$$\text{Cost}(h_\theta(x), y) = \begin{cases} -\log(h_\theta(x)) & \text{if } y = 1 \\ -\log(1 - h_\theta(x)) & \text{if } y = 0 \end{cases}$$

$$\text{cost}(h_\theta(x), y) = -y \log(h_\theta(x)) - (1 - y) \log(1 - h_\theta(x))$$




Optimization

How to minimize the cost function



$$0 \quad \text{cost}(h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1-y) \log(1 - h_{\theta}(x))$$

$$\text{Repeat } \{ \theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \}$$



[Tensorflow Code]

```
def grad(hypothesis, labels):  
    with tf.GradientTape() as tape:  
        loss_value = loss_fn(hypothesis, labels)  
        return tape.gradient(loss_value, [W, b])  
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.01)  
optimizer.apply_gradients(grads_and_vars=zip(grads, [W, b]))
```